



**BITS Pilani**  
Hyderabad Campus

# Database Design & Applications (SS ZG 518)

Dr.R.Gururaj  
CS&IS Dept.

# Transaction Processing(Ch.17)



## **Content**

- ☐ *What is Transaction Model*
- ☐ *Significance of Transaction Model*
- ☐ *States of a transaction*
- ☐ *ACID Properties*
- ☐ *Concurrent Transactions*
- ☐ *Transaction Schedule*
- ☐ *Serial and Concurrent Schedules*
- ☐ *Need for Concurrency Control*
- ☐ *Conflicting Operations*
- ☐ *Conflict Equivalent Schedule & Test for Conflict Serializability*

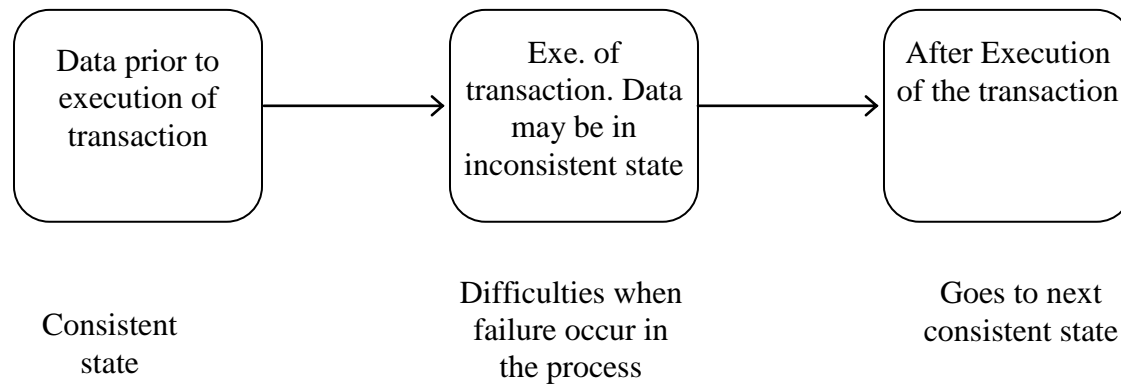
# Transaction Model

A *transaction* is an atomic unit of work that is either completed in its entirety or not done at all.

- For recovery purposes, the system needs to keep track of when the transaction starts, terminates, and commits or aborts.

A *transaction* is a program unit that access and update several data items.

*Read ( )* and *Write ( )* are the basic operations.



Hence, as a result of failure, state of the system will not reflect the state of the real world that the database is supposed to capture.

We call that state as *inconsistent state*.

It is important to define transactions such that they preserve consistency.

## Railway Reservation

Reserve  $x$  number of berths

1. If  $x \leq A$
2.  $A = A - x$
3. Make payment for  $x$  berths
4. Print tickets
5. If  $x > A$  "Then give error message  
"required berths not available"



Database

Ex

if  $x \leq A$

$A = A - x$

make payment

print ticket - Done

else print message

"Required berths  
not available"

User1

$x = 5$

$x \leq A$  True


$A = 10 - 5 = 5$

X  
system crash

User2

$x = 10$

$x \leq A$  False



This is because the berths (5 no) tried by User1 are not given to him but changes are done to database that is why U2 is not able to get them.

# ACID Properties of a Transaction

---

Transaction should possess the following properties called as **ACID** properties.

**Atomicity**: A transaction is an atomic unit of processing. It is either performed in its entirety or not performed at all.

**Consistency Preservation**: The successful execution of a transaction takes the database from one consistent state to another.

**Isolation**: A transaction should be executed as if it is not interfered by any other transaction.

**Durability**: The changes applied to the data by a transaction must be permanent.

# Transaction States



Active State: Initial state when a transaction starts.

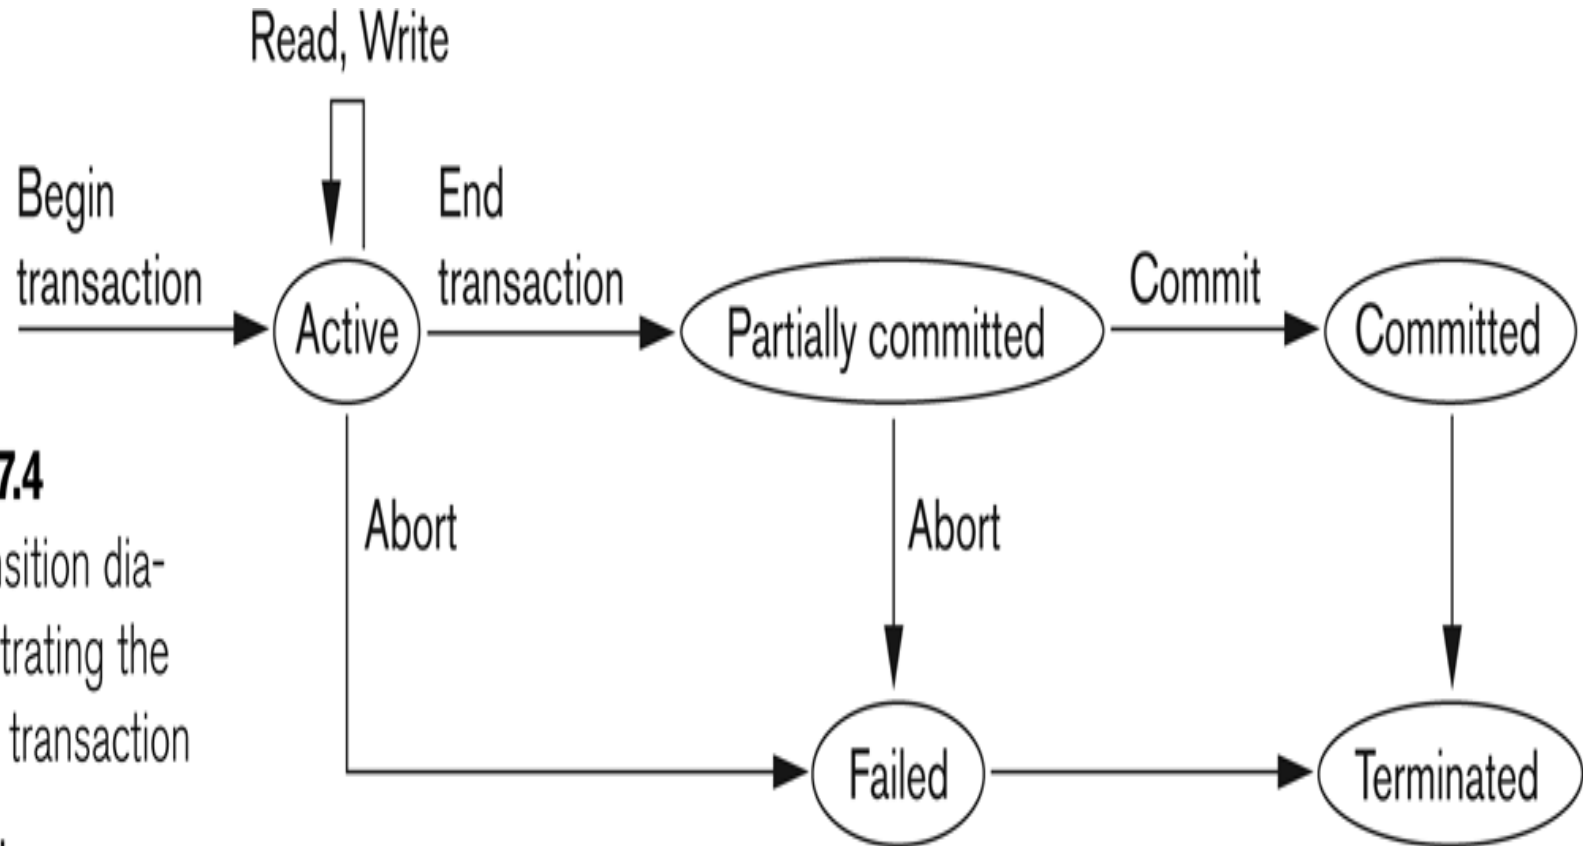
Partially committed State: This state is reached when the last statement is executed and the outcome is not written to the DB.

Failed State: After discovering that the normal execution cannot be continued a transaction is aborted and reaches failed state.

Comitted State: Is reached after successful completion of the transaction.

Terminated State: Is reached after failure or success of the transaction.





**Figure 17.4**

State transition diagram illustrating the states for transaction execution.

# Transaction Schedule

The descriptions that specify the execution sequence of instructions in a set of transactions are called as *schedules*.

Hence schedule can describe the execution sequence of more than one transaction.

Here, in the above schedule the transactions  $T_1$  &  $T_2$  are executed in a serial manner i.e., first all the instructions of the transaction  $T_1$  are executed, and then the instructions of  $T_2$  are executed. Hence the above schedule is known as *serial schedule*.

$T_1$	$T_2$
Read (A) $A = A + 50$ Read (B) $B = B + A$ Write (B)	Read (B) $B = B + 75$ Write (B)

$T_1$  and  $T_2$  are transaction

Read (A) – Reads data item A

Write (B) – Writes the data item B

---

In a serial schedule, instructions belonging to one single transaction appear together.

A serial schedule does not exploit the concurrency. Hence, it is less efficient.

If the transactions are executed concurrently then the resources can be utilized more efficiently hence more throughput is achieved.

A *serial schedule* always results in correct database state that reflect the real world situations.

When the instructions of different transactions of a schedule are executed in an interleaved manner use call such schedules are called *concurrent schedules*.

This kind of concurrent schedules may result in incorrect database state.

## Problems with Concurrent Schedules

### ❑ **The Lost Update Problem**

This occurs when two transactions that access the same database items have their operations interleaved in a way that makes the value of some database item incorrect.

### ❑ **The Temporary Update (or Dirty Read) Problem**

This occurs when one transaction updates a database item and then the transaction fails for some reason (see Section 17.1.4).

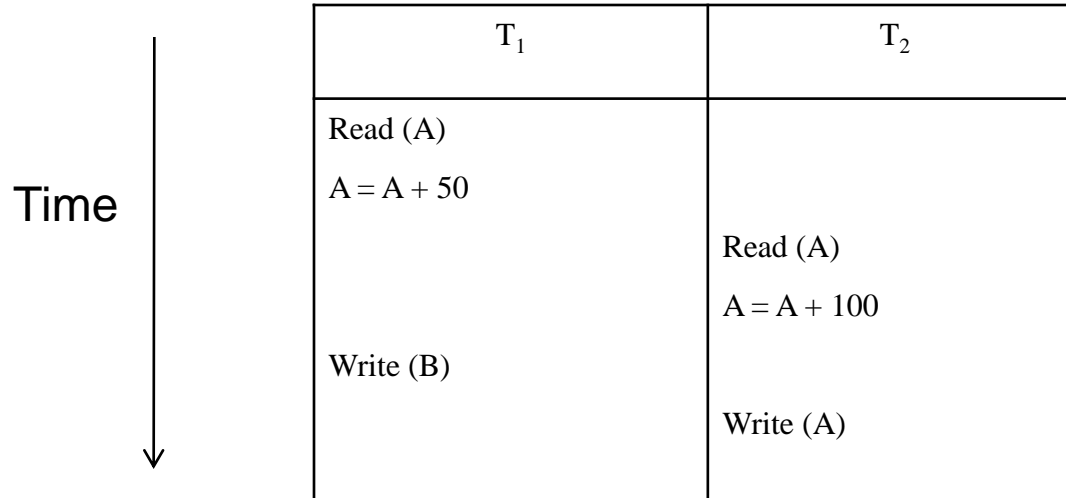
The updated item is accessed by another transaction before it is changed back to its original value.

### ❑ **The Incorrect Summary Problem**

If one transaction is calculating an aggregate summary function on a number of records while other transactions are updating some of these records, the aggregate function may calculate some values before they are updated and others after they are updated.



## Swapping non-conflict pair of operations



Hence it is evident that if we swap non-conflicting operations of a concurrent schedule, it will not affect the final result.

Look at the following example.

T <sub>1</sub>	T <sub>2</sub>
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

(S<sub>1</sub>)

Concurrent schedule  
with T<sub>1</sub> & T<sub>2</sub>  
accessing A, B (data  
item)

T <sub>1</sub>	T <sub>2</sub>
R(A)	
W(A)	
	R(A)
R(B)	
	W(A)
W(B)	
	R(B)
	W(B)

(S<sub>2</sub>)

Swap W(A) in T<sub>2</sub> with  
R(B) in T<sub>1</sub> (because  
they are non  
conflicting)

T <sub>1</sub>	T <sub>2</sub>
R(A)	
W(A)	
R(B)	
	R(A)
W(B)	
	W(A)
	R(B)
	W(B)

(S<sub>3</sub>)

Swap R(A) of T<sub>2</sub> with  
R(B) of T<sub>1</sub> and W(A)  
in T<sub>2</sub> with W(B) in T<sub>1</sub>  
(Since non  
conflicting)



T <sub>1</sub>	T <sub>2</sub>
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
	R(B)
	W(B)

(S<sub>4</sub>)

Swap R(A) of T<sub>2</sub> with  
W(B) of T<sub>1</sub>

Now, the final schedule is a serial schedule

# Conflict Equivalent Schedules



If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are *conflict equivalent*.

Further, we say that a schedule  $S$  is *conflict serializable* if it is conflict equivalent to a serial schedule.

In our example  $S_4$  in the above example is a serial schedule and is conflict equivalent to  $S_1$ . Hence  $S_1$  is a conflict serializable schedule.

$T_1$	$T_2$
R( $\theta$ )	W( $\theta$ )
W( $\theta$ )	

In this schedule we cannot perform any swap between instructions of  $T_1$  and  $T_2$ . Hence it is not conflict serializable

## Test for Conflict Serializability

Let  $S$  be a schedule.

We construct a precedence graph.

Each transaction participating in the schedule will become a vertex.

The set of edges consist of all edges  $T_i \rightarrow T_j$  for which one of the following three conditions hold-


$T_i$  executes  $W(Q)$  before  $T_j$  executes  $R(Q)$

$T_i$  executes  $R(Q)$  before  $T_j$  executes  $W(Q)$

$T_i$  executes  $W(Q)$  before  $T_j$  executes  $W(Q)$

T <sub>0</sub>	T <sub>1</sub>
R(A)	
	R(A)
W(A)	
	W(A)
	R(B)
R(B)	
W(B)	

Infact their schedule is non conflict serializable



T<sub>1</sub> writes(A) after T<sub>0</sub> writes(A) hence we draw on edge from T<sub>0</sub> → T<sub>1</sub>  
 T<sub>1</sub> reads(B) before T<sub>0</sub> write (B) hence we can draw an edge from T<sub>1</sub> → T<sub>0</sub>


At any moment of time, while developing the graph in the above manner, if we see a cycle then the schedule is *not conflict serializable*. If no cycles at the end, then it is *conflict serializable*. Hence the above schedule is not serializable.

Now let us consider the following transaction which is *conflict serializable* and discussed earlier.

T <sub>1</sub>	T <sub>2</sub>
R(A)	R(A)
W(A)	
R(B)	W(A)
W(B)	R(B)
	W(B)

Now, let us draw a precedence graph for the above schedule –

To write A before T<sub>1</sub> reads A. hence we have T<sub>0</sub> → T<sub>1</sub>.



We have only one edge in this graph, and no cycles. Hence it is *conflict serializable*.

## Serial schedule:

- A schedule  $S$  is serial if, for every transaction  $T$  participating in the schedule, all the operations of  $T$  are executed consecutively in the schedule.
- Otherwise, the schedule is called nonserial schedule.

## Serializable schedule:

- A schedule  $S$  is serializable if it is equivalent to some serial schedule of the same  $n$  transactions.

- Being serializable is not the same as being serial
- Being serializable implies that the schedule is a correct schedule.
  - It will leave the database in a consistent state.
  - The interleaving is appropriate and will result in a state as if the transactions were serially executed, yet will achieve efficiency due to concurrent execution.



Ex: Look at the following schedule for the concurrent transactions 1 and 2. The data items are X and Y.

**Schedule** :  $r_2(X)$ ;  $w_2(X)$ ;  $r_1(X)$ ;  $r_2(Y)$ ;  $w_1(X)$ ;  $w_2(Y)$ ;  
 $w_1(X)$ ;  $r_1(Y)$ ;  $w_1(Y)$ ;  $w_2(A)$ ;  $r_1(A)$ ;


Here,  $r_1(X)$  - means that the transaction-1 reads data item X

$w_2(Y)$  - means that the transaction-2 reads data item Y

For the above schedules draw the precedence graph and find if it is conflict serializable.

T <sub>1</sub>	T <sub>2</sub>
	R(x)
	W(x)
R(x)	
	R(y)
W(x)	
	W(y)
W(x)	
R(y)	
W(y)	
	W(x)
R(x)	

Precedence graph



Since there exists a cycle  
it is not serializable.

**Ex:** Look at the following schedules for the concurrent transactions 1 and 2. The data items are A and B.

**Schedule 1 :**  $r_1(A)$ ;  $r_1(B)$ ;  $r_2(A)$ ;  $w_1(B)$ ;  $w_2(A)$ ;  $w_1(B)$ ;  
 $R_2(B)$ ;  $r_1(B)$ ;  $w_2(A)$ ;  $r_1(A)$ ;

**Schedule 2 :**  $r_1(B)$ ;  $r_1(A)$ ;  $r_2(A)$ ;  $w_1(B)$ ;  $w_2(A)$ ;  $R_2(A)$ ;  
 $w_1(B)$ ;  $r_2(B)$ ;  $R_1(B)$ ;  $w_2(A)$ ;

Here,  $r_1(A)$ ; - means that the transaction-1 reads data item A

$w_2(B)$ ; - means that the transaction-2 reads data item B

For each of the above schedules draw the precedence graph and find if they are conflict serializable.

T <sub>1</sub>	T <sub>2</sub>
R(A)	
R(B)	
	R(A)
W(B)	
	W(A)
W(B)	
	R(B)
R(B)	
	W(A)
R(A)	

Schedule 1

T <sub>1</sub>	T <sub>2</sub>
R(B)	
R(A)	
	R(A)
W(B)	
	W(A)
R(A)	
W(B)	
	R(B)
R(B)	
	W(A)

Schedule-2

## **Summary**

- ✓ *What is a transaction*
- ✓ *Properties of a transaction*
- ✓ *States of a transaction*
- ✓ *Transaction execution and the database consistency*
- ✓ *What are concurrent Transactions*
- ✓ *What are Serial and Concurrent Schedules*
- ✓ *Conflict Equivalent Schedule and its importance*
- ✓ *Test for Conflict Serializability*